

In [1]:

```
import os
mingw_path = 'C:\\Program Files\\mingw-w64\\x86_64-5.3.0-posix-seh-rt_v4-rev0\\mingw64\\bin'
os.environ['PATH'] = mingw_path + ';' + os.environ['PATH']
os.environ['PATH'] = mingw_path + ';' + os.environ['PATH']
import xgboost as xgb
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import warnings
from sklearn.cross_validation import train_test_split
from sklearn.linear_model import LinearRegression
warnings.filterwarnings("ignore", category=DeprecationWarning)
from mpl_toolkits.basemap import Basemap
import missingno as msno
import matplotlib.pyplot as pylab
import gmpy2
from pandas.plotting import scatter_matrix
from sklearn.preprocessing import Imputer
import numpy as np
import numpy as np
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
sns.set(style="whitegrid", color_codes=True)
pylab.rcParams['figure.figsize'] = 16, 12
pd.options.mode.chained_assignment = None
%matplotlib inline
```

C:\Users\sugan\Anaconda3\lib\site-packages\sklearn\cross_validation.py:44:
 DeprecationWarning: This module was deprecated in version 0.18 in favor of
 the model_selection module into which all the refactored classes and func-
 tions are moved. Also note that the interface of the new CV iterators are d-
 ifferent from that of this module. This module will be removed in 0.20.
 "This module will be removed in 0.20.", DeprecationWarning)

In [2]:

```
train = pd.read_csv('train_2017.csv') #this data includes log error and parcel ID of year 2017
prop = pd.read_csv('properties_2017.csv') #this data includes log error and parcel ID of year 2017
train.head()
```

C:\Users\sugan\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:2717: DtypeWarning: Columns (49) have mixed types. Specify dtype option on import or set low_memory=False.
interactivity=interactivity, compiler=compiler, result=result)

Out[2]:

	parcelid	logerror	transactiondate
0	14297519	0.025595	2017-01-01
1	17052889	0.055619	2017-01-01
2	14186244	0.005383	2017-01-01
3	12177905	-0.103410	2017-01-01
4	10887214	0.006940	2017-01-01

chaging the data types of the given data

In [3]:

```
categorical = ['airconditioningtypeid', 'architecturalstyletypeid', 'buildingclasstypeid', 'fips', 'hashottuborspa',
               'propertycountylandusecode', 'propertyzoningdesc', 'rawcensustractandblock', 'censustractandblock', 'regionidcounty', 'regionidcity',
               'regionidzip', 'regionidneighborhood', 'buildingqualitytypeid', 'decktypeid', 'heatingorsystemtypeid', 'pooltypeid10', 'pooltypeid2', 'pooltypeid7',
               'propertylandusetypeid', 'storytypeid', 'typeconstructiontypeid', 'taxdelinquencyflag', 'taxdelinquencyyear']
for col in categorical:
    prop[col] = prop[col].astype('category')
train['transactiondate'] = pd.to_datetime(train['transactiondate'])
```

Renaming the columns

In [4]:

```

data_dict = pd.read_excel('zillow_data_dictionary.xlsx', parse_cols="A:B")
features = data_dict['Feature'].apply(lambda x:x.strip(""))
data_dict['Feature'] = features
d = {} #forming a dictionary to rename columns
for a,b in data_dict.iterrows():
    d[b['Feature']] = b['Rename']

#Rename Columns
prop.rename_axis(d,axis=1,inplace=True)

#Making a copy of data frame for future analysis purpose
prop_copy = prop.copy()

```

Merging the house attributes data and house transaction data

In [6]:

```

merge_df = prop.merge(train, how='inner')
total_df = prop.merge(train, how='left')
total_df['transactiondate'] = pd.to_datetime(total_df['transactiondate'])

```

In [8]:

```

merge_df['transactiondate'] = pd.to_datetime(merge_df['transactiondate'])
merge_df['month'] = merge_df['transactiondate'].dt.month
merge_df['day'] = merge_df['transactiondate'].dt.day
merge_df['quarter'] = merge_df['transactiondate'].dt.quarter
merge_df['transaction_year'] = merge_df['transactiondate'].dt.year
merge_df['age'] = 2017 - merge_df['yearbuilt']
merge_df['month'] = merge_df['month'].astype('category')
merge_df['quarter'] = merge_df['quarter'].astype('category')

```

View of the data


In [9]:

```
merge_df.head()
```

Out[9]:

	parcelid	ac_type	arch_Type	basementsqft	total_bathcnt	bedroomcnt	building_
0	17054981	NaN	NaN	NaN	5.0	4.0	NaN
1	17055743	NaN	NaN	NaN	2.0	3.0	NaN
2	17068109	NaN	NaN	NaN	1.5	3.0	NaN
3	17073952	NaN	NaN	NaN	2.0	2.0	NaN
4	17078502	NaN	NaN	NaN	1.0	2.0	NaN

5 rows × 65 columns



- Data consists of 64 columns and we can see that Null values are present we have to process the data to remove them

In [10]:

```
merge_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 77613 entries, 0 to 77612
Data columns (total 65 columns):
parcelid                77613 non-null int64
ac_type                 25007 non-null category
arch_Type               207 non-null category
basementsqft           50 non-null float64
total_bathcnt           77579 non-null float64
bedroomcnt              77579 non-null float64
building_class          15 non-null category
building_quality        49809 non-null category
calculatedbathnbr       76963 non-null float64
decktypeid              614 non-null category
firstfloor_finisharea   6037 non-null float64
total_finisharea        77378 non-null float64
finished_living         73923 non-null float64
perimeter_living        42 non-null float64
total_area              3027 non-null float64
firstfloor_finisharea   6037 non-null float64
basetotalarea           386 non-null float64
fips                    77579 non-null category
fireplacecnt            8289 non-null float64
fullbathcnt             76963 non-null float64
garagecarcnt            25520 non-null float64
garagetotalsqft         25520 non-null float64
hashottuborspa          1539 non-null category
heatingid               49571 non-null category
latitude                 77579 non-null float64
longitude                77579 non-null float64
lotsizesquarefeet       69321 non-null float64
poolcnt                 16174 non-null float64
poolsizesum             869 non-null float64
pooltypeid10            465 non-null category
pooltypeid2             1074 non-null category
pooltypeid7             15079 non-null category
landusecode             77579 non-null category
landusetypeid           77579 non-null category
zoningdesc              50476 non-null category
rawcensustractandblock  77579 non-null category
city                    76107 non-null category
county                  77579 non-null category
neighborhood            30974 non-null category
zip                     77529 non-null category
roomcnt                 77579 non-null float64
storytypeid             50 non-null category
3/4bathnbr              10106 non-null float64
typeconstructiontypeid  223 non-null category
unitcnt                 50703 non-null float64
yardbuildingsqft17      2393 non-null float64
yardbuildingsqft26      70 non-null float64
yearbuilt               77309 non-null float64
numberofstories         17599 non-null float64
fireplaceflag           172 non-null object
structuretaxvaluedollarcnt 77464 non-null float64
totaltax                77578 non-null float64
assessmentyear          77579 non-null float64
landtaxvaluedollarcnt   77577 non-null float64
taxperyear              77574 non-null float64
taxdelinquencyflag      2900 non-null category
taxdelinquencyyear      2900 non-null category
censustractandblock     77332 non-null category

```

logerror77613non-nullfloat64transactiondate77613non-nulldatetime64[ns]month77613non-nullcategoryday77613non-nullint64quarter77613non-nullcategorytransaction_year77613non-nullint64age77309non-nullfloat64dtypes: category(26), datetime64[ns](1), float64(34), int64(3), object(1)memory usage: 33.2+ MB

Descriptive Statistics (count, mean, std, min, Max) of numerical attributes in the the data

In [11]:

merge_df.describe()

Out[11]:

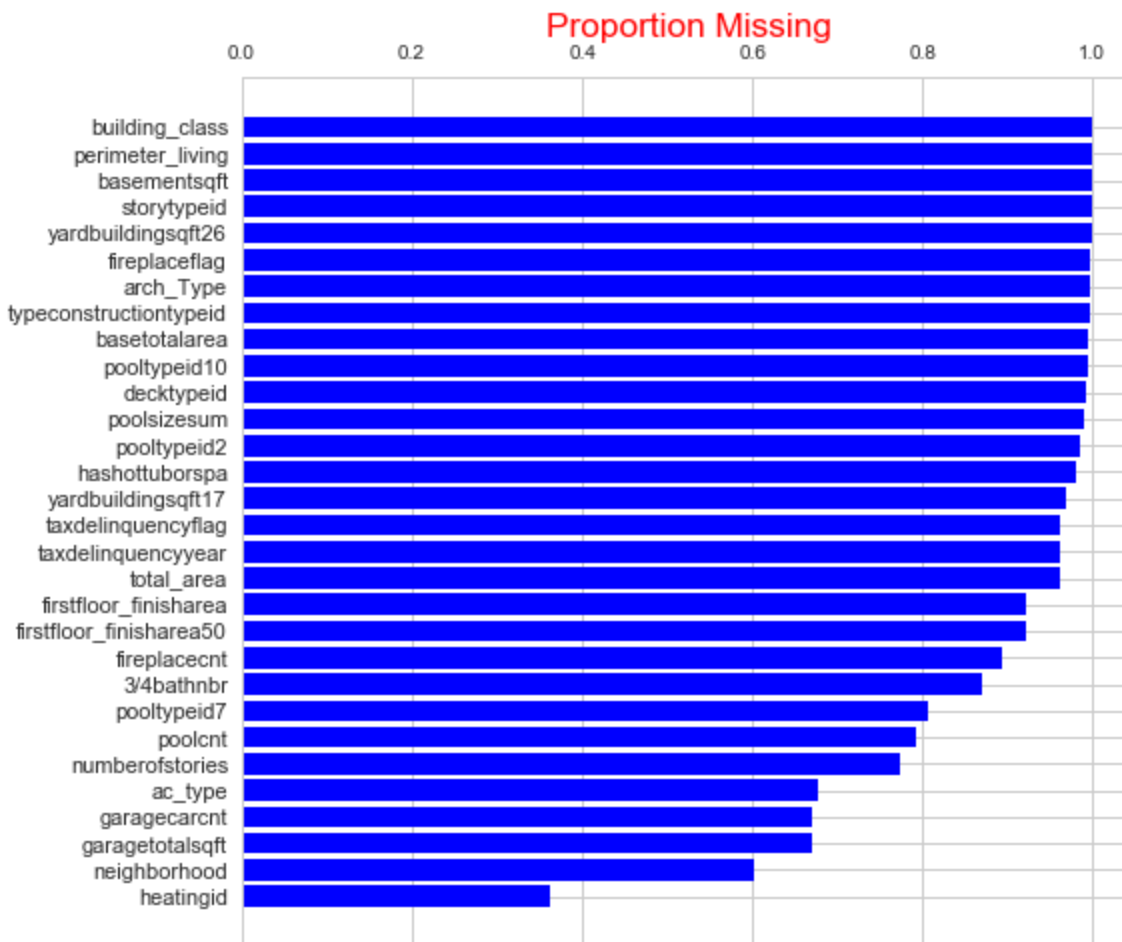
	parcelid	basementsqft	total_bathcnt	bedroomcnt	calculatedbathnbr
count	7.761300e+04	50.000000	77579.000000	77579.000000	76963.000000
mean	1.300781e+07	679.720000	2.298496	3.053223	2.316392
std	3.518717e+06	689.703546	0.996732	1.140480	0.979689
min	1.071186e+07	38.000000	0.000000	0.000000	1.000000
25%	1.153821e+07	273.000000	2.000000	2.000000	2.000000
50%	1.253004e+07	515.000000	2.000000	3.000000	2.000000
75%	1.421101e+07	796.500000	3.000000	4.000000	3.000000
max	1.676893e+08	3560.000000	18.000000	16.000000	18.000000

8 rows × 37 columns

Percentage and Graphical analysis of Null values in the data


In [104]:

```
total_nulls = merge_df.isnull().sum().sort_values(ascending = False)
percent = ((merge_df.isnull().sum().sort_values())/(len(merge_df))).sort_values(ascending = False)
missin_val = pd.concat([total_nulls, percent], axis = 1, keys = ['Total', 'Percent_missing'])
missin_val.sort_values('Percent_missing', ascending = False)
missin_val.reset_index(inplace=True)
missin_val.rename({'index': 'Column_name'})
missin_val.sort_values(by= ['Percent_missing'], axis = 0, ascending = True, inplace=True)
missin_val.head()
plt.figure(figsize=(8,8))
plt.barh(bottom= np.arange(30), width= missin_val.Percent_missing.values[35:], tick_label = missin_val['index'].values[35:], color = 'blue')
plt.xlabel('Proportion Missing', fontsize = 'xx-large', color = 'red')
plt.tick_params(axis = 'y', labelsize = 11.0)
ax = plt.gca()
ax.xaxis.tick_top()
ax.xaxis.set_label_position('top')
plt.savefig('missing_values')
```



visualizing the Null values in a dataframe Graphically


```
merge_df = merge_df.reindex_axis(sorted(merge_df.columns), axis=1)
msno.matrix(merge_df, figsize= (24, 15), labels = True, width ratios=(15,1))
```



```
## For Exploratory data analysis i am not including the log error values which are below 1 percent and above 99 percent
nomerge_df = merge_df[(merge_df['logerror'] < np.percentile(merge_df['logerror'], 99))
    & (merge_df['logerror'] > np.percentile(merge_df['logerror'], 1))]
omerge_df = merge_df[~((merge_df['logerror'] < np.percentile(merge_df['logerror'], 99))
    & (merge_df['logerror'] > np.percentile(merge_df['logerror'], 1)))]
```


Zestimate (log error) analysis

In [13]:

```
plt.figure(figsize = (10,8))  
ax1 = sns.distplot(nomerge_df.logerror.values)  
ax1.set_xlabel('logerror', size =15, color = 'red')
```

Out[13]:

<matplotlib.text.Text at 0x1a6ff7464e0>



In [14]:

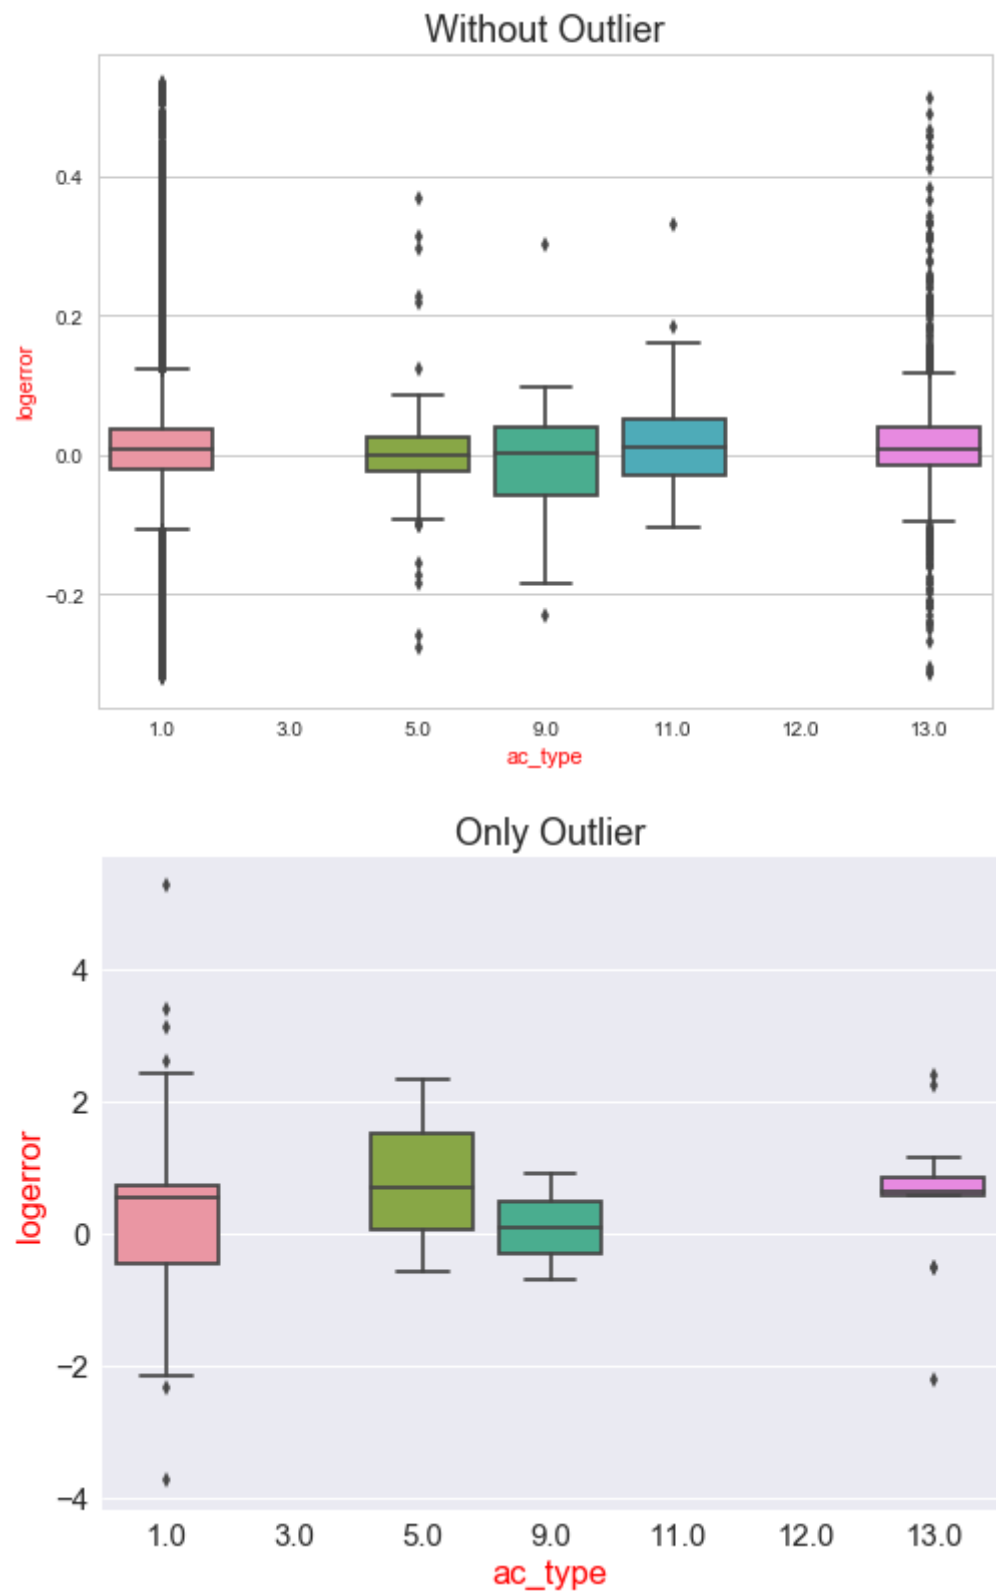
```
merge_df['abs_logerror']= merge_df.logerror.abs()  
nomerge_df['abs_logerror']= nomerge_df.logerror.abs()  
omerge_df['abs_logerror']= merge_df.logerror.abs()
```

In [15]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['ac_type'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('Without Outlier')

# -----
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = omerge_df['ac_type'], y= omerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('Only Outlier')
```

```
Out[15]:  
<matplotlib.text.Text at 0x1a6f272d940>
```

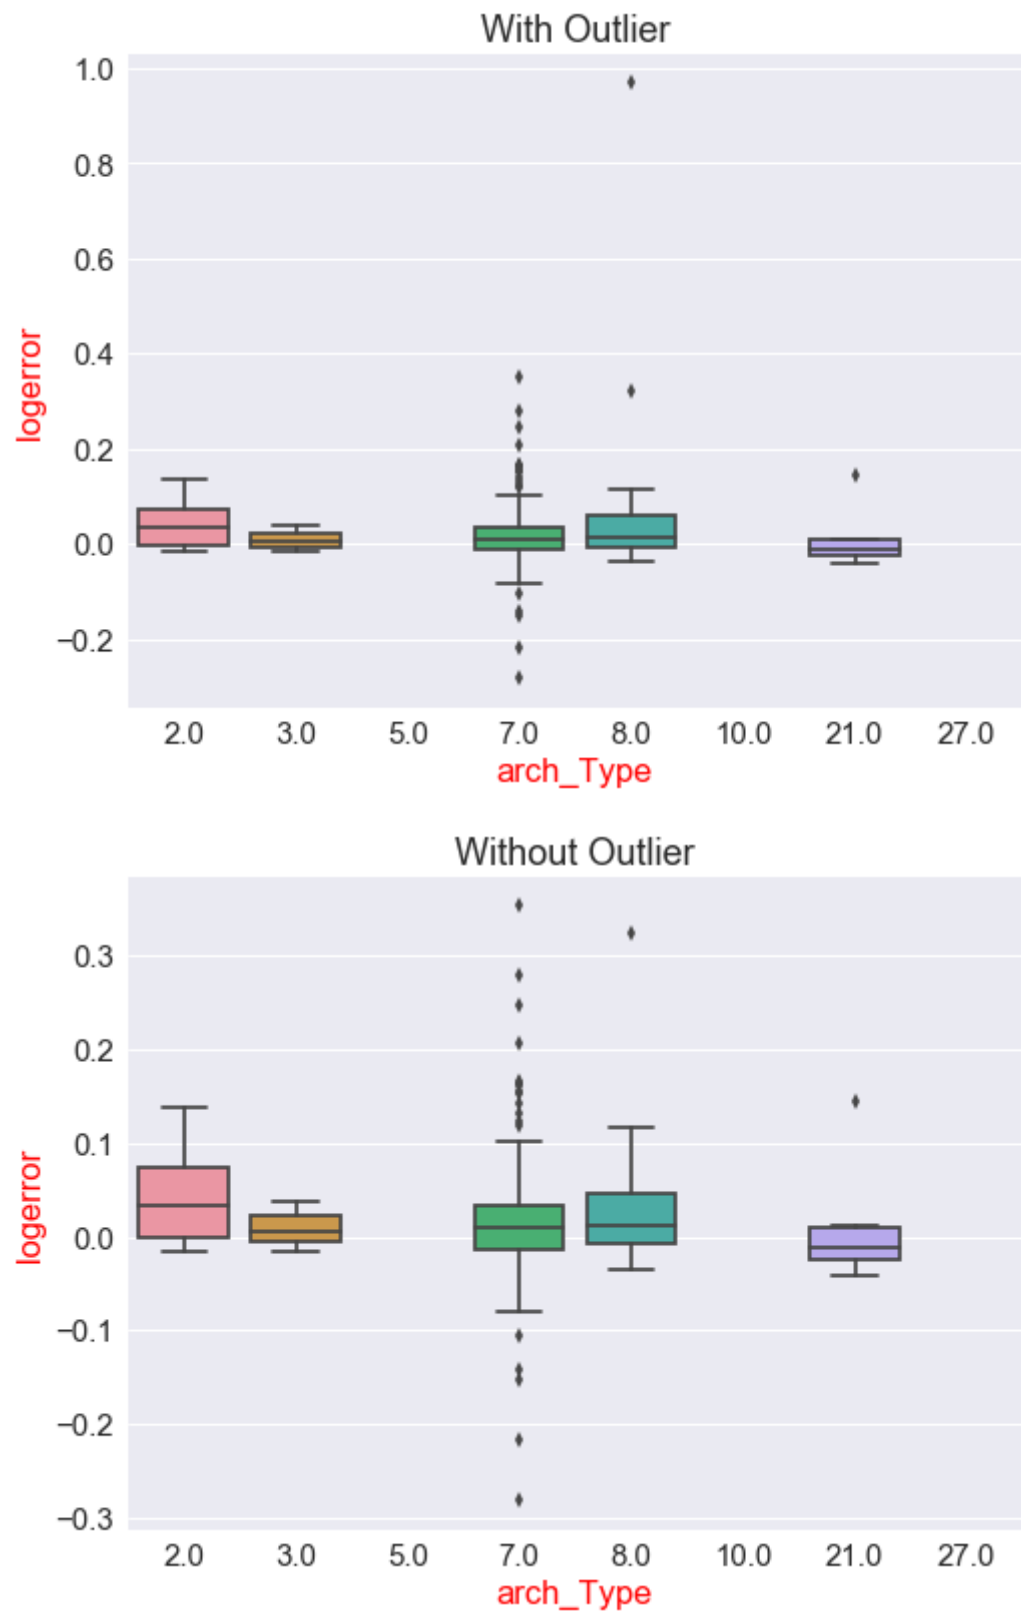


In [111]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = merge_df['arch_Type'], y= merge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('With Outlier')

#-----
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['arch_Type'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('Without Outlier')
```

```
Out[111]:  
<matplotlib.text.Text at 0x1d7fc7e9390>
```



In [112]:

```
categorical_cols = merge_df.dtypes[merge_df.dtypes == 'category'].index.tolist()
categorical_cols
```

Out[112]:

```
['ac_type',
 'arch_Type',
 'building_class',
 'building_quality',
 'censustractandblock',
 'city',
 'county',
 'decktypeid',
 'fips',
 'hashottuborspa',
 'heatingid',
 'landusecode',
 'landusetypeid',
 'month',
 'neighborhood',
 'pooltypeid10',
 'pooltypeid2',
 'pooltypeid7',
 'quarter',
 'rawcensustractandblock',
 'storytypeid',
 'taxdelinquencyflag',
 'taxdelinquencyyear',
 'typeconstructiontypeid',
 'zip',
 'zoningdesc']
```


In [113]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = merge_df['building_quality'], y= merge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('With Outlier')

#-----
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['building_quality'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('Without Outlier')
```



```
Out[113]:  
<matplotlib.text.Text at 0x1d7f9ae4d68>
```

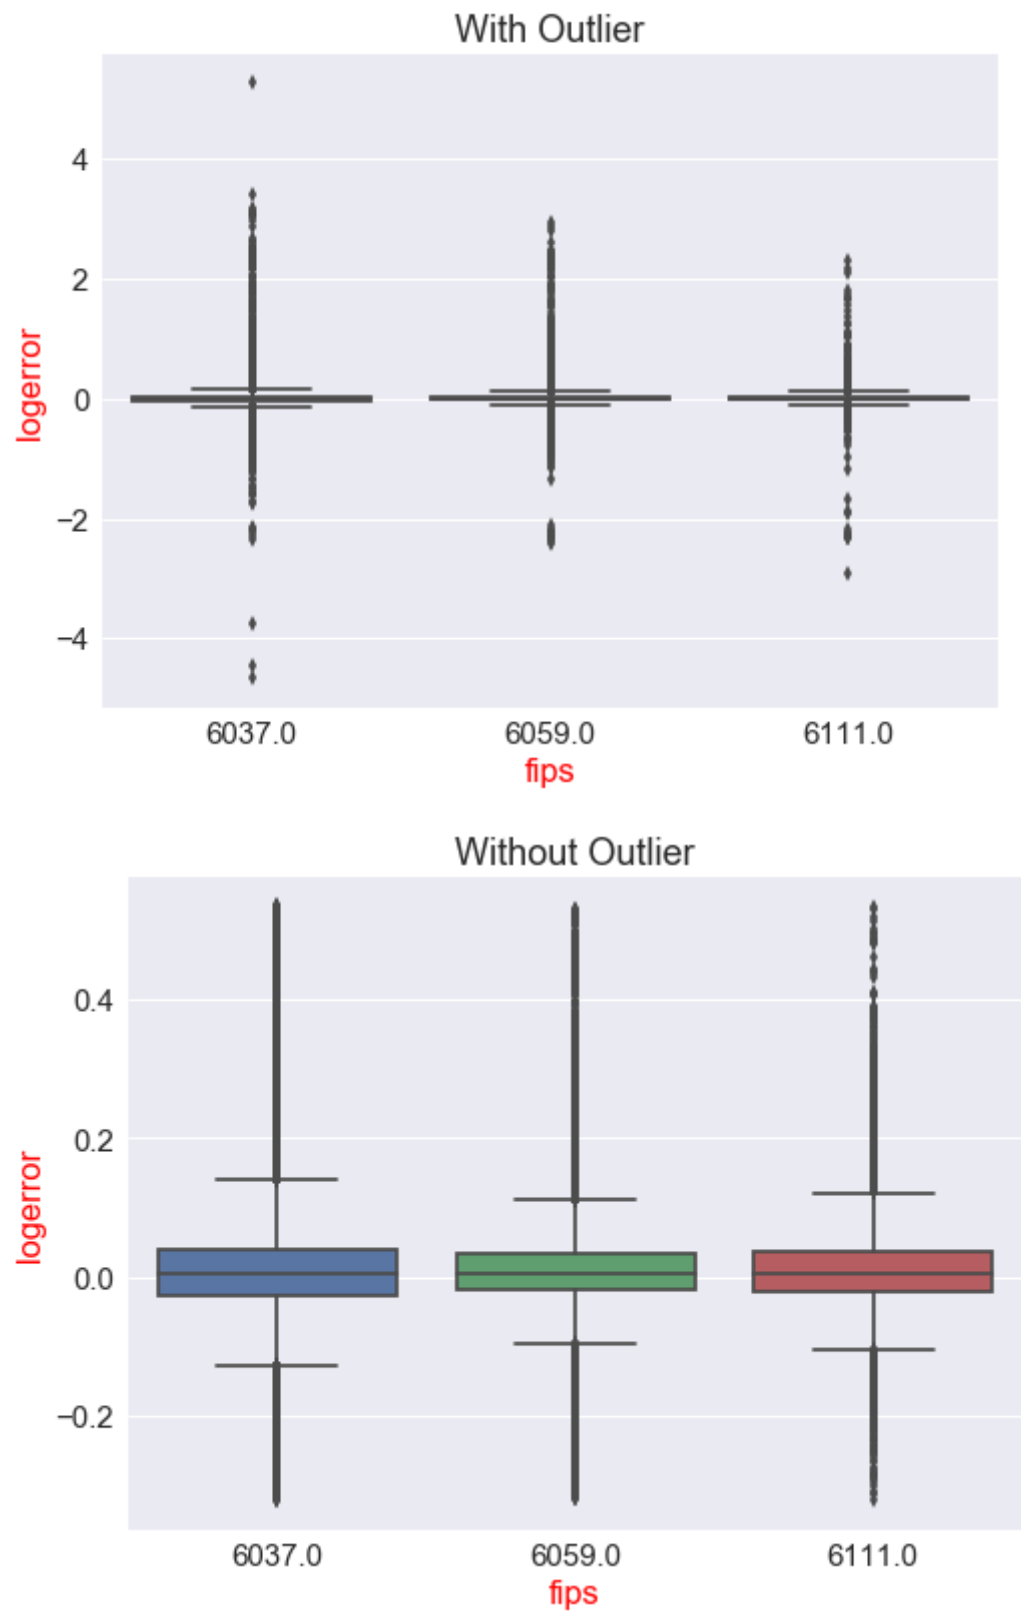


In [114]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = merge_df['fips'], y= merge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('With Outlier')

#-----
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['fips'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('Without Outlier')
```

```
Out[114]:  
<matplotlib.text.Text at 0x1d7f9494080>
```

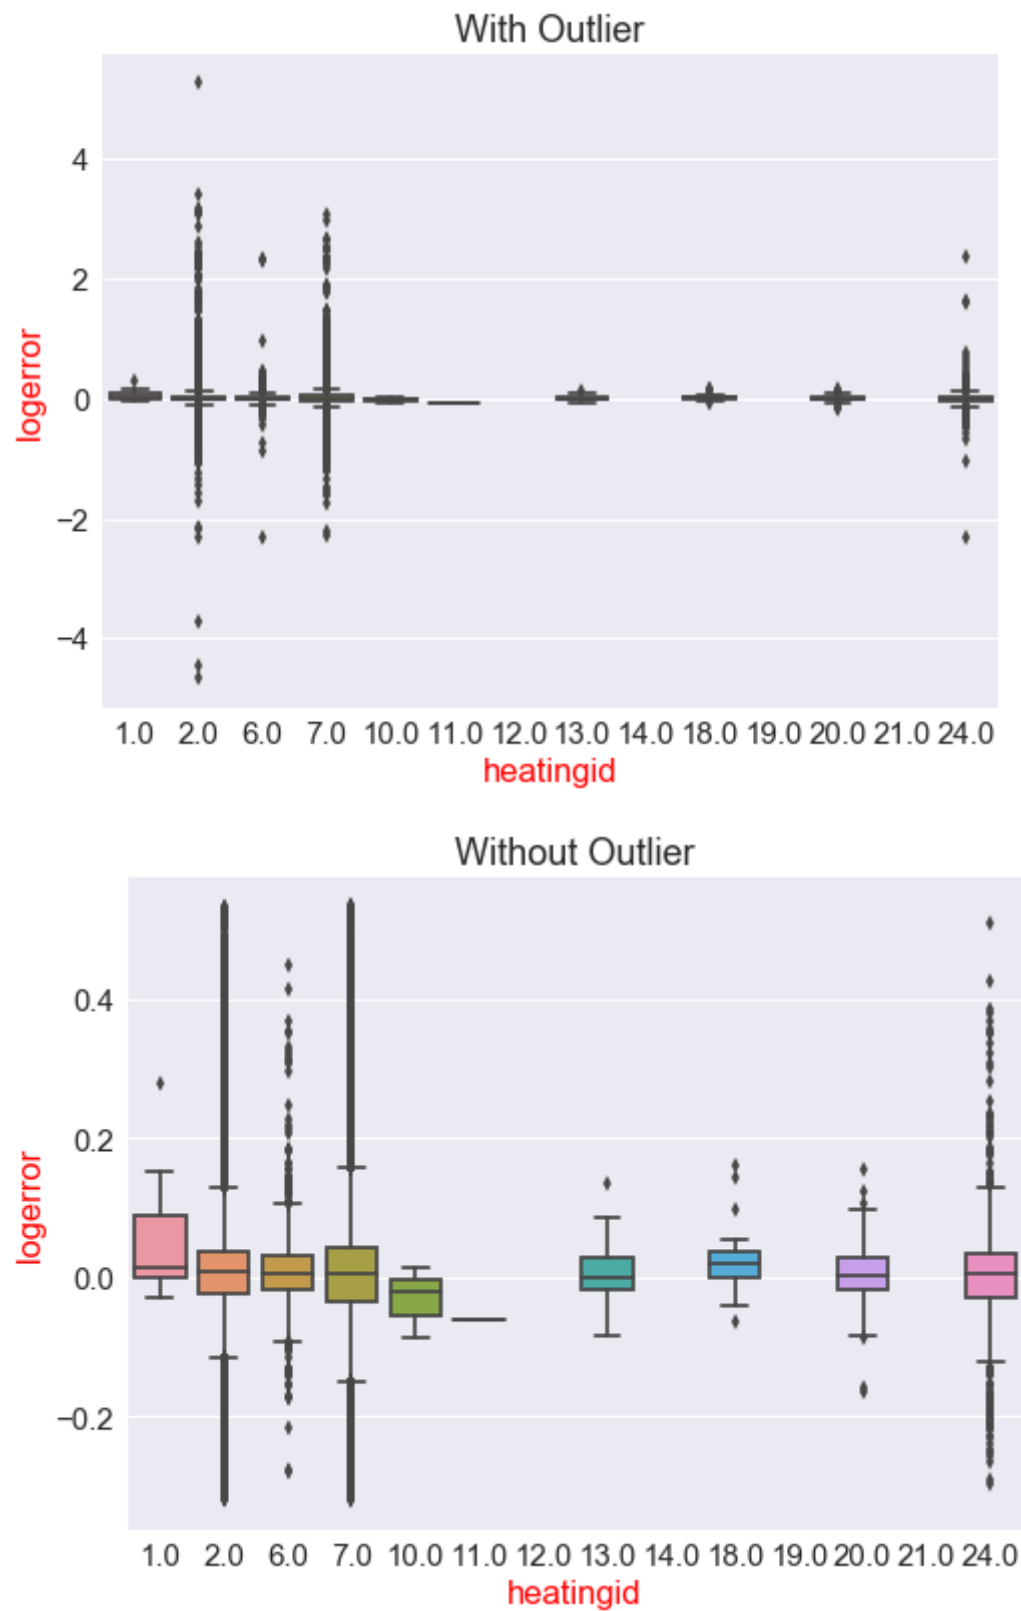


In [115]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = merge_df['heatingid'], y= merge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('With Outlier')

#-----
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['heatingid'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('Without Outlier')
```

```
Out[115]:  
<matplotlib.text.Text at 0x1d7fc947320>
```





In [116]:


```
bools = merge_df.groupby(by='landusetypeid').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(merge_df.groupby(by='landusetypeid').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(merge_df.groupby(by='landusetypeid').mean()['abs_logerror'][bools].
index.tolist(), size = 'small')
ax1.set_xlabel('landusetypeid', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('With Outliers', fontsize = 16)

#-----
----
bools = nomerge_df.groupby(by='landusetypeid').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(nomerge_df.groupby(by='landusetypeid').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(nomerge_df.groupby(by='landusetypeid').mean()['abs_logerror'][bools]
.index.tolist(), size = 'small')
ax1.set_xlabel('landusetypeid', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('Without Outliers', fontsize = 16)

#-----
-----
bools = omerge_df.groupby(by='landusetypeid').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(omerge_df.groupby(by='landusetypeid').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(omerge_df.groupby(by='landusetypeid').mean()['abs_logerror'][bools]
.index.tolist(), size = 'small')
ax1.set_xlabel('landusetypeid', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('only Outliers', fontsize = 16)
```

```
Out[116]:  
<matplotlib.text.Text at 0x1d7e099b668>
```





In [117]:

```

bools = merge_df.groupby(by='zoningdesc').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(merge_df.groupby(by='zoningdesc').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(merge_df.groupby(by='zoningdesc').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('zoningdesc', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('With Outliers', fontsize = 16)
plt.title('Only Outliers', fontsize = 16)
plt.xticks(rotation=45)



#-----
----
bools = nomerge_df.groupby(by='zoningdesc').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(nomerge_df.groupby(by='zoningdesc').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(nomerge_df.groupby(by='zoningdesc').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('zoningdesc', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.xticks(rotation=45)


#-----
-----
bools = omerge_df.groupby(by='zoningdesc').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(omerge_df.groupby(by='zoningdesc').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(omerge_df.groupby(by='zoningdesc').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('zoningdesc', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('Outliers', fontsize = 16)
plt.xticks(rotation=70)

```

Out[117]:

```
(array([ -50.,   0.,   50.,  100.,  150.,  200.,  250.,  300.,  350.]),  
<a list of 9 Text xticklabel objects>)
```





In [118]:

```

bools = merge_df.groupby(by='city').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(merge_df.groupby(by='city').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(merge_df.groupby(by='city').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('city', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('With Outliers', fontsize = 16)
plt.xticks(rotation=45)



#-----
----
bools = nomerge_df.groupby(by='city').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(nomerge_df.groupby(by='city').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(nomerge_df.groupby(by='city').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('city', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('No Outliers', fontsize = 16)
plt.xticks(rotation=45)


#-----
-----
bools = omerge_df.groupby(by='city').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(omerge_df.groupby(by='city').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(omerge_df.groupby(by='city').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('city', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('Only Outliers', fontsize = 16)
plt.xticks(rotation=70)

```

Out[118]:

```
(array([ -20.,   0.,  20.,  40.,  60.,  80., 100., 120., 140., 160.]),  
<a list of 10 Text xticklabel objects>)
```





In [119]:

```

bools = merge_df.groupby(by='zip').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(merge_df.groupby(by='zip').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(merge_df.groupby(by='zip').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('zip', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('With Outliers', fontsize = 16)
plt.title('Only Outliers', fontsize = 16)
plt.xticks(rotation=45)


#-----
----
bools = nomerge_df.groupby(by='zip').mean()['logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(nomerge_df.groupby(by='zip').mean()['logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(nomerge_df.groupby(by='zip').mean()['logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('zip', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.xticks(rotation=45)


#-----
-----
bools = omerge_df.groupby(by='zip').mean()['abs_logerror'].notnull()
plt.figure(figsize = (8,6))
plt.plot(omerge_df.groupby(by='zip').mean()['abs_logerror'][bools].values)
ax1 = plt.gca()
ax1.set_xticklabels(omerge_df.groupby(by='zip').mean()['abs_logerror'][bools].index.tolist(), size = 'small')
ax1.set_xlabel('zip', size = 'small', color = 'red')
ax1.set_ylabel('mean_logerror', size = 'small', color = 'red')
plt.title('Outliers', fontsize = 16)
plt.xticks(rotation=70)

```

Out[119]:

```
(array([ -50.,   0.,   50.,  100.,  150.,  200.,  250.,  300.,  350.,  400.]),  
<a list of 10 Text xticklabel objects>)
```






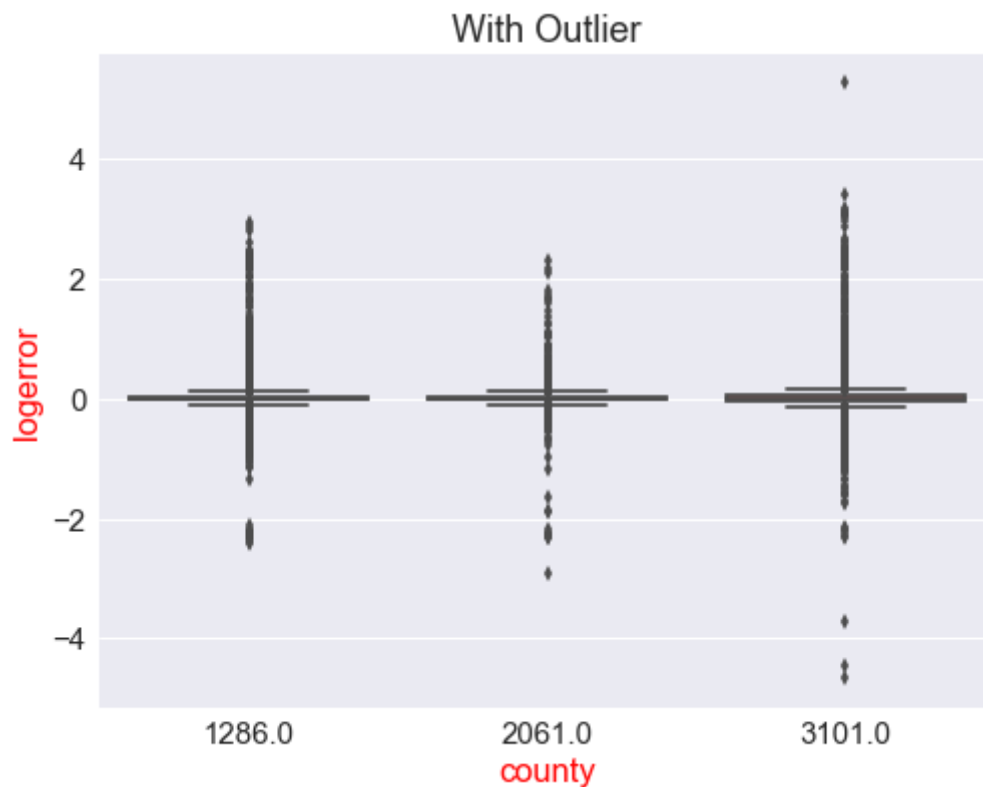
In [120]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = merge_df['county'], y= merge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('With Outlier')

#-----
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['county'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color ='red')
g.set_title('Without Outlier')
```

Out[120]:

<matplotlib.text.Text at 0x1d7f7a56b70>



from the above analysis unnecessary category attributes are removed

- if they have large number of null values (90% of values being null)
- no deviation in the mean log error with the category and standar deviation being same
- attributes which give same amount of information (for examples, city and zip code provide insights to same information)

In [16]:

```
copy_merge_df = merge_df.copy() #making a copy of dataframe before altering it  
copy_nomerge_df = nomerge_df.copy()  
copy_omerge_df = omerge_df.copy()
```

In [17]:

```
copy_merge_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 77613 entries, 0 to 77612
Data columns (total 66 columns):
parcelid                77613 non-null int64
ac_type                 25007 non-null category
arch_Type               207 non-null category
basementsqft           50 non-null float64
total_bathcnt           77579 non-null float64
bedroomcnt              77579 non-null float64
building_class          15 non-null category
building_quality        49809 non-null category
calculatedbathnbr       76963 non-null float64
decktypeid              614 non-null category
firstfloor_finisharea   6037 non-null float64
total_finisharea        77378 non-null float64
finished_living         73923 non-null float64
perimeter_living        42 non-null float64
total_area              3027 non-null float64
firstfloor_finisharea   6037 non-null float64
basetotalarea           386 non-null float64
fips                    77579 non-null category
fireplacecnt            8289 non-null float64
fullbathcnt             76963 non-null float64
garagecarcnt            25520 non-null float64
garagetotalsqft         25520 non-null float64
hashottuborspa          1539 non-null category
heatingid               49571 non-null category
latitude                77579 non-null float64
longitude                77579 non-null float64
lotsizesquarefeet       69321 non-null float64
poolcnt                 16174 non-null float64
poolsizesum             869 non-null float64
pooltypeid10            465 non-null category
pooltypeid2             1074 non-null category
pooltypeid7             15079 non-null category
landusecode             77579 non-null category
landusetypeid           77579 non-null category
zoningdesc              50476 non-null category
rawcensustractandblock  77579 non-null category
city                    76107 non-null category
county                  77579 non-null category
neighborhood            30974 non-null category
zip                     77529 non-null category
roomcnt                 77579 non-null float64
storytypeid             50 non-null category
3/4bathnbr              10106 non-null float64
typeconstructiontypeid  223 non-null category
unitcnt                 50703 non-null float64
yardbuildingsqft17      2393 non-null float64
yardbuildingsqft26      70 non-null float64
yearbuilt               77309 non-null float64
numberofstories         17599 non-null float64
fireplaceflag           172 non-null object
structuretaxvaluedollarcnt 77464 non-null float64
totaltax                77578 non-null float64
assessmentyear          77579 non-null float64
landtaxvaluedollarcnt   77577 non-null float64
taxperyear              77574 non-null float64
taxdelinquencyflag      2900 non-null category
taxdelinquencyyear      2900 non-null category
censustractandblock     77332 non-null category

```

```

logerror          77613 non-null float64
transactiondate   77613 non-null datetime64[ns]
month             77613 non-null category
day              77613 non-null int64
quarter          77613 non-null category
transaction_year  77613 non-null int64
age              77309 non-null float64
abs_logerror      77613 non-null float64
dtypes: category(26), datetime64[ns](1), float64(35), int64(3), object
(1)
memory usage: 33.8+ MB

```

In [18]:

```

merge_df.drop(['landusecode', 'building_class', 'decktypeid', 'fips', 'pooltypeid10',
'pooltypeid2', 'pooltypeid7', 'county',
'typeconstructiontypeid', 'storytypeid', 'taxdelinquencyflag', 'taxdelinq
uencyyear', 'zip', 'neighborhood'], axis = 1, inplace =True)
nmerge_df.drop(['landusecode', 'building_class', 'decktypeid', 'fips', 'pooltypeid10',
'pooltypeid2', 'pooltypeid7', 'county',
'typeconstructiontypeid', 'storytypeid', 'taxdelinquencyflag', 'taxdelinq
uencyyear', 'zip', 'neighborhood'], axis = 1, inplace =True)
omerge_df.drop(['landusecode', 'building_class', 'decktypeid', 'fips', 'pooltypeid10',
'pooltypeid2', 'pooltypeid7', 'county',
'typeconstructiontypeid', 'storytypeid', 'taxdelinquencyflag', 'taxdelinq
uencyyear', 'zip', 'neighborhood'], axis = 1, inplace =True)

```

Numerical attribute analysis

In [19]:

```

numerical_col = merge_df.dtypes[(merge_df.dtypes == 'int64')|(merge_df.dtypes == 'float6
4')].index.tolist()

```

In [20]:

```
merge_df[numerical_col].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 77613 entries, 0 to 77612
Data columns (total 40 columns):
parcelid                77613 non-null int64
basementsqft            50 non-null float64
total_bathcnt           77579 non-null float64
bedroomcnt              77579 non-null float64
calculatedbathnbr       76963 non-null float64
firstfloor_finisharea   6037 non-null float64
firstfloor_finisharea   6037 non-null float64
total_finisharea        77378 non-null float64
finished_living         73923 non-null float64
perimeter_living        42 non-null float64
total_area              3027 non-null float64
firstfloor_finisharea   6037 non-null float64
firstfloor_finisharea   6037 non-null float64
basetotalarea           386 non-null float64
fireplacecnt            8289 non-null float64
fullbathcnt             76963 non-null float64
garagecarcnt            25520 non-null float64
garagetotalsqft         25520 non-null float64
latitude                77579 non-null float64
longitude               77579 non-null float64
lotsizesquarefeet       69321 non-null float64
poolcnt                 16174 non-null float64
poolsizesum             869 non-null float64
roomcnt                 77579 non-null float64
3/4bathnbr              10106 non-null float64
unitcnt                 50703 non-null float64
yardbuildingsqft17      2393 non-null float64
yardbuildingsqft26      70 non-null float64
yearbuilt               77309 non-null float64
numberofstories         17599 non-null float64
structuretaxvaluedollarcnt 77464 non-null float64
totaltax                77578 non-null float64
assessmentyear          77579 non-null float64
landtaxvaluedollarcnt   77577 non-null float64
taxperyear              77574 non-null float64
logerror                77613 non-null float64
day                     77613 non-null int64
transaction_year        77613 non-null int64
age                     77309 non-null float64
abs_logerror            77613 non-null float64
dtypes: float64(37), int64(3)
memory usage: 26.8 MB
```


In [21]:

```
numerical_col
```

Out[21]:

```
['parcelid',  
 'basementsqft',  
 'total_bathcnt',  
 'bedroomcnt',  
 'calculatedbathnbr',  
 'firstfloor_finisharea',  
 'total_finisharea',  
 'finished_living',  
 'perimeter_living',  
 'total_area',  
 'firstfloor_finisharea',  
 'basetotalarea',  
 'fireplacecnt',  
 'fullbathcnt',  
 'garagecarcnt',  
 'garagetotalsqft',  
 'latitude',  
 'longitude',  
 'lotsizesquarefeet',  
 'poolcnt',  
 'poolsizesum',  
 'roomcnt',  
 '3/4bathnbr',  
 'unitcnt',  
 'yardbuildingsqft17',  
 'yardbuildingsqft26',  
 'yearbuilt',  
 'numberofstories',  
 'structuretaxvaluedollarcnt',  
 'totaltax',  
 'assessmentyear',  
 'landtaxvaluedollarcnt',  
 'taxperyear',  
 'logerror',  
 'day',  
 'transaction_year',  
 'age',  
 'abs_logerror']
```

In [127]:

```
corr_mat = merge_df.corr()
plt.figure(figsize = (30, 20))
sns.set(font_scale = 1.25)
sns.heatmap(corr_mat, annot = True)
```


Out[127]:

<matplotlib.axes._subplots.AxesSubplot at 0x1d7f7a1b240>



In [128]:

```
#coorelation matrix arranged in descending order of correlation with log error
cols = corr_mat.nlargest(15, 'logerror')['logerror'].index
plt.figure(figsize= (20,20))
coor_mat_10 = merge_df[cols].corr()
sns.heatmap(coor_mat_10, square = True, annot=True, linewidths =0.5)
sns.set_style(style = 'darkgrid')
sns.set(font_scale = 1)
```




In [129]:

```
plt.figure(figsize=(8, 6))
g = sns.factorplot(x = 'month', y= 'abs_logerror', data= merge_df, estimator= np.mean,
size = 6, aspect=2)
sns.set(font_scale=1.5)
g.set_axis_labels("month", "mean_abs_logerror")
g.set_xticklabels(labels =['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep'])
g.set_xlabel(color = 'red')
g.set_ylabel(color = 'red')
```

Out[129]:

<seaborn.axisgrid.FacetGrid at 0x1d7f7b04128>

<matplotlib.figure.Figure at 0x1d7f1de7208>




In [130]:

```
plt.figure(figsize=(8, 6))
g = sns.factorplot(x = 'yearbuilt', y = 'abs_logerror', data= merge_df, estimator= np.me
an, size = 6, aspect=2, ci = None)
sns.set(font_scale=1.5)
g.set_axis_labels("year_built", "mean_abs_logerror")
g.set_xticklabels(labels =[])
g.set_xlabel(color = 'red')
g.set_ylabel(color = 'red')
```

Out[130]:

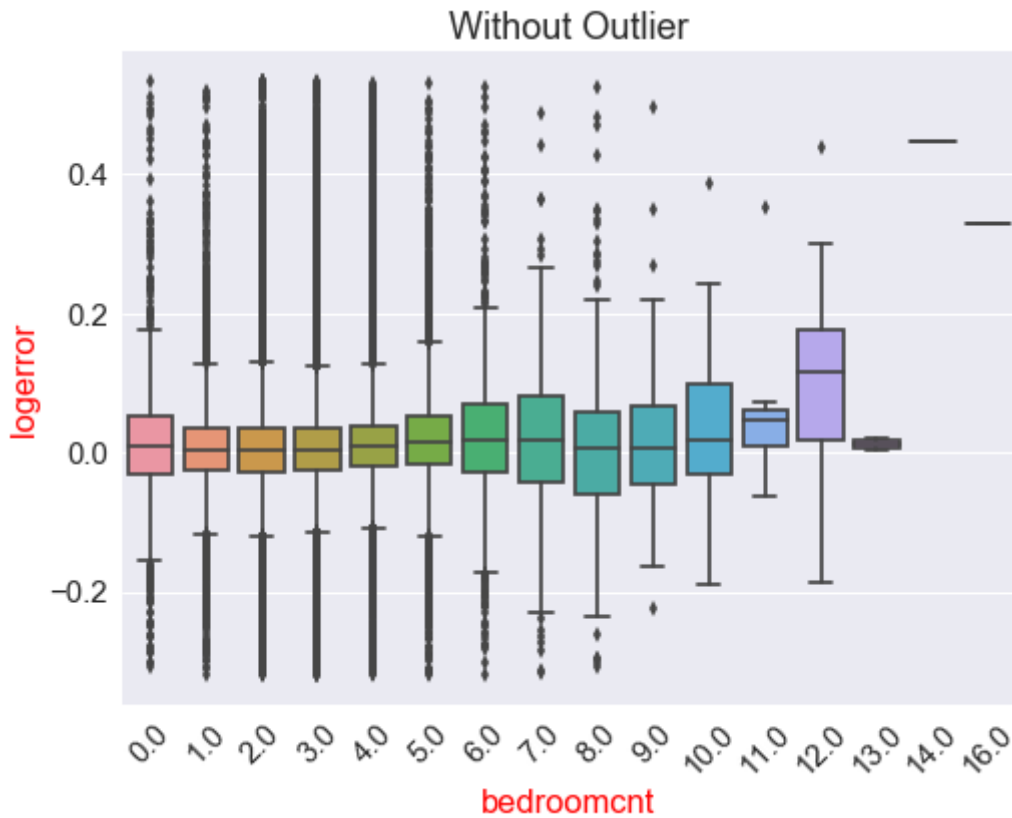
<seaborn.axisgrid.FacetGrid at 0x1d77fa57ef0>

<matplotlib.figure.Figure at 0x1d7efa9cb70>




In [131]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['bedroomcnt'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color='red')
g.set_title('Without Outlier')
for item in g.get_xticklabels():
    item.set_rotation(45)
```




In [132]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['fullbathcnt'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color = 'red')
g.set_title('Without Outlier')
for item in g.get_xticklabels():
    item.set_rotation(45)
```



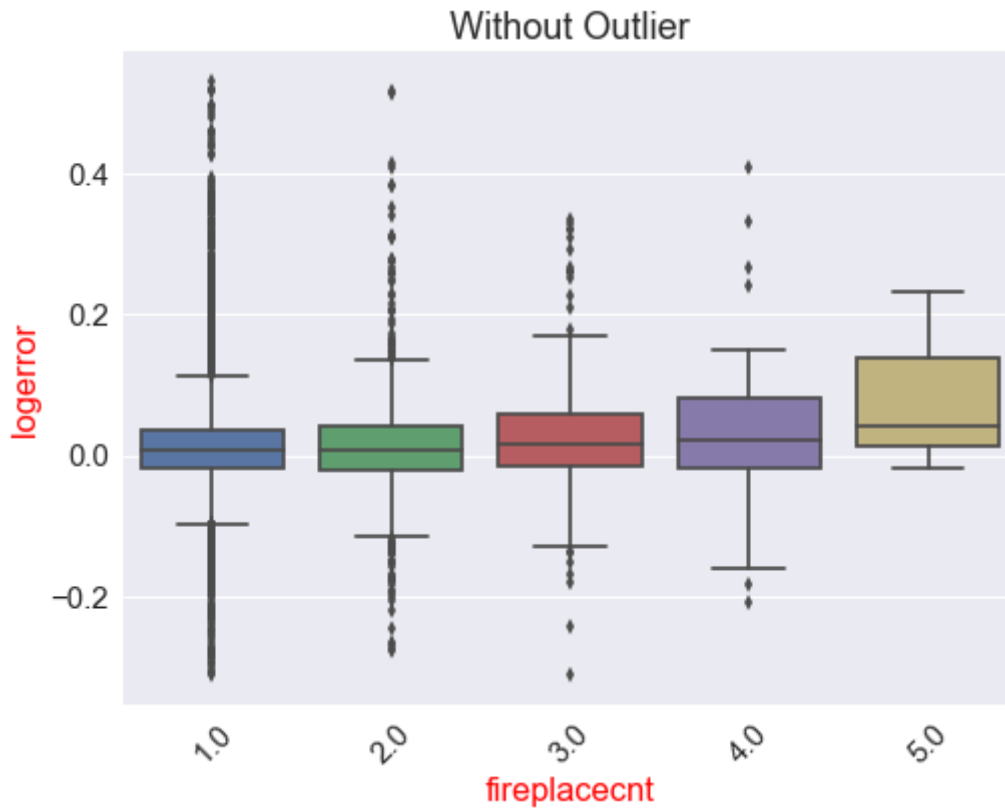
In [133]:

```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['garagecarcnt'], y= nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color = 'red')
g.set_title('Without Outlier')
for item in g.get_xticklabels():
    item.set_rotation(45)
```





In [134]:



```
plt.figure(figsize=(8, 6))
g = sns.boxplot(x = nomerge_df['fireplacecnt'], y = nomerge_df['logerror'])
sns.set(font_scale=1.5)
g.set_xlabel(g.get_xlabel(), color = 'red')
g.set_ylabel(g.get_ylabel(), color = 'red')
g.set_title('Without Outlier')
for item in g.get_xticklabels():
    item.set_rotation(45)
```



In [135]:

```
for feature in ['total_finisharea', 'garagetotalsqft', 'lotsizesquarefeet', 'totaltax']:  
    a = sns.jointplot(merge_df[feature], merge_df['logerror'], size = 7, space = 0.1)  
    a.set_axis_labels(xlabel= feature, ylabel='logerror', color = 'red', size = 15)
```






In [136]:

```
scatter_matrix(merge_df[['total_finisharea', 'garagetotalsqft', 'lotsizesquarefeet', 'totaltax', 'logerror']], figsize = (15,10))
```

Out[136]:

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000001D7E0147C8
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D7F1D2686
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D722A2C2E
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D723334F2
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72339990
8>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x000001D72339994
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D7238E7B7
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72C3864E
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72C41616
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72C4266D
8>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x000001D72EC8EBE
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72ECED90
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F1DD6A
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F239F6
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F2BB27
8>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F31D20
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F359B3
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F3E87B
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F3F6DA
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F4B819
8>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F505F6
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F57DC1
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F5E35F
8>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F65D8D
0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x000001D72F6BE86
0>]], dtype=object)
```



dropping the columns with more than 68 percent of missing values and imputing the missing values with mean in the case of numerical variable and most frequent in the case of categorical variable

In [22]:

```
cp_merge_df = merge_df.copy()
cp_nomerge_df = nomerge_df.copy()
cp_omerge_df = omerge_df.copy()
```

In [23]:

```
merge_df.dropna(thresh=0.327*len(merge_df), axis=1, inplace=True)
nomerge_df.dropna(thresh=0.327*len(nomerge_df), axis=1, inplace=True)
omerge_df.dropna(thresh=0.327*len(omerge_df), axis=1, inplace=True)
```

#removing the rows which have null values (the row is removed even if it has one null value)

impute the missing values with mean for numerical variables and mostfrequent values for categorical variables


In [24]:

```
imp_df =merge_df.copy()
```

```
for col in imp_df.columns:
    if imp_df[col].dtype == 'float64' and col not in ('month', 'day', 'quarter'):
        imp_df[col].fillna(float(imp_df[col].mean()), inplace = True)
    if imp_df[col].dtype == 'int64' and col not in ('month', 'day', 'quarter'):
        imp_df[col].fillna(int(imp_df[col].mean()), inplace = True)
for col in imp_df.columns:
    if imp_df[col].dtype not in ('float64', 'int64'):
        imp_df[col] = imp_df[col].astype(dtype= object)
        imp_df[col].fillna(imp_df[col].mode().values[0], inplace = True)
        imp_df[col] = imp_df[col].astype('category')
```


```
imp_mat = imp_df.corr()
plt.figure(figsize = (30, 20))
sns.set(font_scale = 1.25)
sns.heatmap(imp_mat, annot = True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x1d72ee9a710>
```



In [142]:

```
cols = imp_mat.nlargest(15, 'logerror')['logerror'].index
plt.figure(figsize= (20,20))
imp_mat_15 = merge_df[cols].corr()
sns.heatmap(imp_mat_15, square = True, annot=True, linewidths =0.5)
sns.set_style(style = 'darkgrid')
sns.set(font_scale = 1)
```



In [26]:

```
from sklearn.preprocessing import LabelEncoder
```

Gradient Boosting Variable Importance

In [144]:

```
y = imp_df['logerror']
x = imp_df.drop(['logerror', 'abs_logerror', 'yearbuilt', 'day', 'month', 'quarter', 'transactiondate', 'parcelid'], axis=1)

for c in x.columns:
    if (x[c].dtype.name == 'category'):
        le = LabelEncoder()
        le.fit(x[c].values)
        x[c]=le.transform(x[c].values)
```

In [145]:

```
xgb_params = {
    'eta': 0.05,
    'max_depth': 8,
    'subsample': 0.7,
    'colsample_bytree': 0.7,
    'objective': 'reg:linear',
    'silent': 1,
    'seed': 0,
    'lambda': 5
}


train = xgb.DMatrix(x,y)
boost_model = xgb.train(xgb_params, train, num_boost_round=150)
```

In [146]:

```
fig, ax = plt.subplots(1,1,figsize= (20, 13))
xgb.plot_importance(boost_model, grid = False, height= 0.8, ax=ax)
```

Out[146]:

<matplotlib.axes._subplots.AxesSubplot at 0x1d7f1e3ee80>



feature engineering with the help of gradient boosting and correlation matrix

- finished_living and total_finisharea have same correlation so total_finisharea is picked
- bedroomcnt, calculatebathnbr, full_bathcnt, total_bathcnt have same correlation so bedroomcnt and total bath cnt are merged to total no of living rooms
- taxperyear, structuretax value dollar cnt, land tax value dollar count have same correlation so total tax is picked

In [27]:

```
imp_df['totalroomcnt'] = imp_df['total_bathcnt'] + imp_df['bedroomcnt']
```

In [148]:

```
def funct(c):  
    if (c['logerror'] > np.percentile(imp_df.logerror.values, 95)):  
        val = 'positive_outlier'  
    elif (c['logerror'] < np.percentile(imp_df.logerror.values, 5)):  
        val = 'negative_outlier'  
    else:  
        val = 'not_outlier'  
    return val  
imp_df['log_group'] = imp_df.apply(funct, axis=1)
```

Variable used for regression

- 'age', 'totalroomcnt', 'total_finisharea', 'latitude', 'longitude', 'totaltax', 'lotsizesquarefeet'

In [28]:

```
imp_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 77613 entries, 0 to 77612
Data columns (total 36 columns):
parcelid                77613 non-null int64
total_bathcnt           77613 non-null float64
bedroomcnt              77613 non-null float64
building_quality        77613 non-null category
calculatedbathnbr       77613 non-null float64
total_finisharea        77613 non-null float64
finished_living         77613 non-null float64
fullbathcnt             77613 non-null float64
garagecarcnt            77613 non-null float64
garagetotalsqft         77613 non-null float64
heatingid               77613 non-null category
latitude                77613 non-null float64
longitude               77613 non-null float64
lotsizesquarefeet       77613 non-null float64
landusetypeid           77613 non-null category
zoningdesc              77613 non-null category
rawcensustractandblock  77613 non-null category
city                    77613 non-null category
roomcnt                 77613 non-null float64
unitcnt                 77613 non-null float64
yearbuilt               77613 non-null float64
structuretaxvaluedollarcnt 77613 non-null float64
totaltax                77613 non-null float64
assessmentyear          77613 non-null float64
landtaxvaluedollarcnt   77613 non-null float64
taxperyear              77613 non-null float64
censustractandblock     77613 non-null category
logerror                77613 non-null float64
transactiondate         77613 non-null category
month                   77613 non-null category
day                     77613 non-null int64
quarter                 77613 non-null category
transaction_year        77613 non-null int64
age                     77613 non-null float64
abs_logerror            77613 non-null float64
totalroomcnt            77613 non-null float64
dtypes: category(10), float64(23), int64(3)
memory usage: 20.6 MB
```

Linear Regression

In [30]:

```
imp_df[['age', 'totalroomcnt', 'total_finisharea', 'latitude', 'longitude', 'totaltax',
'lotsizesquarefeet' ]].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 77613 entries, 0 to 77612
Data columns (total 7 columns):
age                77613 non-null float64
totalroomcnt       77613 non-null float64
total_finisharea   77613 non-null float64
latitude           77613 non-null float64
longitude          77613 non-null float64
totaltax           77613 non-null float64
lotsizesquarefeet  77613 non-null float64
dtypes: float64(7)
memory usage: 4.7 MB
```

In [31]:

```
X = imp_df[['age', 'totalroomcnt', 'total_finisharea', 'latitude', 'longitude', 'totaltax',
'lotsizesquarefeet']]
Y = imp_df['logerror']

#n = pd.get_dummies(imp_df.log_group)
#X = pd.concat([X, n], axis=1)
#m = pd.get_dummies(imp_df.censustractandblock)
#X = pd.concat([X, m], axis=1)
#drops = ['censustractandblock']
#X.drop(drops, inplace=True, axis=1)
X.head()
```

Out[31]:

	age	totalroomcnt	total_finisharea	latitude	longitude	totaltax	lotsizesqu
0	35.0	9.0	3760.0	34449407.0	-119254052.0	872850.0	42688.0000
1	66.0	5.0	1444.0	34454169.0	-119237898.0	436157.0	7108.0000
2	38.0	4.5	1698.0	34365693.0	-119448392.0	286606.0	2588.0000
3	28.0	4.0	986.0	34305600.0	-119284000.0	258888.0	29973.4370
4	69.0	3.0	1170.0	34278012.0	-119257047.0	592930.0	5643.0000

In [37]:

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.4, random_state=0)
regressor = LinearRegression(normalize=False)
regressor.fit(X_train, y_train,)
```

Out[37]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

In [154]:

```
y_pred = regressor.predict(X_test)
#print('Linear Regression R squared: ', regressor.score(X_test, y_test))
lin_mse = mean_squared_error(y_pred, y_test)
lin_rmse = np.sqrt(lin_mse)
print('Linear Regression RMSE: ', lin_rmse)
print('Linear Regression AME: ', mean_absolute_error(y_pred, y_test))
```

Linear Regression RMSE: 0.173057429591

Linear Regression AME: 0.0708969525796

In [155]:

```
print('Coefficients: \n', regressor.coef_)
```

Coefficients:


```
[ 1.44328811e-05 -1.60302881e-03  1.26306182e-05 -4.42832689e-09
 8.21271980e-09 -6.20066147e-09  1.84672058e-08]
```

In [156]:

```
plt.scatter(y_pred, (y_pred- y_test))
```

Out[156]:

<matplotlib.collections.PathCollection at 0x1d7bfc1e668>




In [157]:

```
sns.distplot((y_pred- y_test),norm_hist=False)
```

Out[157]:

<matplotlib.axes._subplots.AxesSubplot at 0x1d734823e80>



Ridge Regression

In [158]:

```
from sklearn.linear_model import Ridge, RidgeCV, ElasticNet, LassoCV, LassoLarsCV
from sklearn.model_selection import cross_val_score
```

In [159]:

```
RMSE = []
alphas = [0.05, 0.1, 0.3, 1, 3, 5, 10, 15, 30, 50, 75, 100, 125, 150, 175, 200, 225, 250, 1000]
for tuning_param in alphas:
    ridge_model = Ridge(alpha=tuning_param)
    ridge_model.fit(X_train, y_train)
    ridge_model.score(X_test, y_test)
    print('Coefficients: ', tuning_param, '\n', ridge_model.coef_)
    #print(cross_val_score(ridge_model, X_train, y_train, scoring=None, cv = 5))
    RMSE.append((sum((ridge_model.predict(X_test) - y_test)**2)/len(y_test))**0.5)
```


Coefficients: 0.05

[1.44328861e-05	-1.60302764e-03	1.26306163e-05	-4.42832743e-09
	8.21271963e-09	-6.20066116e-09	1.84672065e-08]	

Coefficients: 0.1

[1.44328912e-05	-1.60302646e-03	1.26306143e-05	-4.42832797e-09
	8.21271946e-09	-6.20066084e-09	1.84672072e-08]	

Coefficients: 0.3

[1.44329113e-05	-1.60302176e-03	1.26306065e-05	-4.42833013e-09
	8.21271877e-09	-6.20065959e-09	1.84672098e-08]	

Coefficients: 1

[1.44329819e-05	-1.60300531e-03	1.26305792e-05	-4.42833770e-09
	8.21271638e-09	-6.20065521e-09	1.84672189e-08]	

Coefficients: 3

[1.44331835e-05	-1.60295830e-03	1.26305012e-05	-4.42835934e-09
	8.21270955e-09	-6.20064268e-09	1.84672451e-08]	

Coefficients: 5

[1.44333852e-05	-1.60291129e-03	1.26304231e-05	-4.42838097e-09
	8.21270271e-09	-6.20063016e-09	1.84672712e-08]	

Coefficients: 10

[1.44338892e-05	-1.60279378e-03	1.26302281e-05	-4.42843503e-09
	8.21268563e-09	-6.20059885e-09	1.84673366e-08]	

Coefficients: 15

[1.44343931e-05	-1.60267630e-03	1.26300331e-05	-4.42848909e-09
	8.21266855e-09	-6.20056754e-09	1.84674020e-08]	

Coefficients: 30

[1.44359044e-05	-1.60232394e-03	1.26294483e-05	-4.42865123e-09
	8.21261732e-09	-6.20047366e-09	1.84675980e-08]	

Coefficients: 50

[1.44379185e-05	-1.60185437e-03	1.26286689e-05	-4.42886730e-09
	8.21254905e-09	-6.20034855e-09	1.84678593e-08]	

Coefficients: 75

[1.44404344e-05	-1.60126779e-03	1.26276954e-05	-4.42913720e-09
	8.21246377e-09	-6.20019226e-09	1.84681856e-08]	

Coefficients: 100

[1.44429485e-05	-1.60068164e-03	1.26267225e-05	-4.42940691e-09
	8.21237855e-09	-6.20003609e-09	1.84685118e-08]	

Coefficients: 125

[1.44454607e-05	-1.60009593e-03	1.26257504e-05	-4.42967642e-09
	8.21229339e-09	-6.19988003e-09	1.84688377e-08]	

Coefficients: 150

[1.44479711e-05	-1.59951064e-03	1.26247790e-05	-4.42994574e-09
	8.21220829e-09	-6.19972409e-09	1.84691633e-08]	

Coefficients: 175

[1.44504796e-05	-1.59892577e-03	1.26238083e-05	-4.43021485e-09
	8.21212326e-09	-6.19956826e-09	1.84694887e-08]	

Coefficients: 200

[1.44529862e-05	-1.59834134e-03	1.26228383e-05	-4.43048377e-09
	8.21203829e-09	-6.19941254e-09	1.84698139e-08]	

Coefficients: 225

[1.44554910e-05	-1.59775733e-03	1.26218690e-05	-4.43075250e-09
	8.21195338e-09	-6.19925694e-09	1.84701388e-08]	

Coefficients: 250

[1.44579940e-05	-1.59717375e-03	1.26209004e-05	-4.43102103e-09
	8.21186853e-09	-6.19910145e-09	1.84704635e-08]	

Coefficients: 1000


[1.45322344e-05	-1.57986244e-03	1.25921683e-05	-4.43898667e-09
	8.20935157e-09	-6.19448903e-09	1.84800949e-08]	

In [160]:

```
plt.plot(alphas, RMSE)
```

Out[160]:

[<matplotlib.lines.Line2D at 0x1d72ca887b8>]



Random Forest

In [161]:

```
from sklearn.ensemble import RandomForestRegressor
forest_reg = RandomForestRegressor(random_state=42)
forest_reg.fit(X_train, y_train)
```

Out[161]:

```
RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                       max_features='auto', max_leaf_nodes=None,
                       min_impurity_split=1e-07, min_samples_leaf=1,
                       min_samples_split=2, min_weight_fraction_leaf=0.0,
                       n_estimators=10, n_jobs=1, oob_score=False, random_state=42,
                       verbose=0, warm_start=False)
```

In [162]:

```
y_pred = forest_reg.predict(X_test)
forest_mse = mean_squared_error(y_pred, y_test)
forest_rmse = np.sqrt(forest_mse)
print('Random Forest RMSE: ', forest_rmse)
print('Random Forest AME: ', mean_absolute_error(y_pred, y_test))
```


```
Random Forest RMSE: 0.184491342476
Random Forest AME: 0.0862187977687
```

In [163]:

```
plt.scatter(y_pred, (y_pred- y_test))
```

Out[163]:

<matplotlib.collections.PathCollection at 0x1d72fa799e8>



Gradient Boosting

In [38]:

```
from sklearn import ensemble
from sklearn.ensemble import GradientBoostingRegressor
model = ensemble.GradientBoostingRegressor(n_estimators=100, max_depth=4, min_samples_s
plit=40, learning_rate=0.01)
model.fit(X_train, y_train)
```

Out[38]:

```
GradientBoostingRegressor(alpha=0.9, criterion='friedman_mse', init=None,
learning_rate=0.01, loss='ls', max_depth=4, max_features=Non
e,
max_leaf_nodes=None, min_impurity_split=1e-07,
min_samples_leaf=1, min_samples_split=40,
min_weight_fraction_leaf=0.0, n_estimators=100,
presort='auto', random_state=None, subsample=1.0, verbose=0,
warm_start=False)
```

In [165]:

```
y_pred = model.predict(X_test)
model_mse = mean_squared_error(y_pred, y_test)
model_rmse = np.sqrt(model_mse)
#print('Gradient Boosting R squared": %.4f' % model.score(X_test, y_test))
print('Gradient Boosting RMSE: ', model_rmse)
print('Gradient Boosting RMSE: ', mean_absolute_error(y_pred, y_test))
```

Gradient Boosting RMSE: 0.173087825491


Gradient Boosting RMSE: 0.0709582747472

In [166]:

```
plt.scatter(y_pred, (y_pred- y_test))
```

Out[166]:

<matplotlib.collections.PathCollection at 0x1d7311f2cc0>



Cross_validation

In [51]:

```
from sklearn.cross_validation import cross_val_score
#Features after Cross Validation
features=['age', 'totalroomcnt', 'total_finisharea', 'latitude']
x = X[features]
gb_score=cross_val_score(model,x,Y,cv=10,scoring='mean_absolute_error')
-gb_score.mean()
```

Out[51]:

0.070531605180707557

In [52]:

```
X1_train, X1_test, Y1_train, Y1_test = train_test_split(x, Y, test_size=0.4, random_state=0)
```

In [53]:

```
#XGB after Cross validation
model = ensemble.GradientBoostingRegressor(n_estimators=100, max_depth=4, min_samples_s
plit=40, learning_rate=0.01)
model.fit(X1_train, Y1_train)
y_pred = model.predict(X1_test)
model_mse = mean_squared_error(y_pred, y_test)
#print('Gradient Boosting R squared": %.4f' % model.score(X_test, y_test))
print('Gradient Boosting MAE: ', mean_absolute_error(y_pred, y_test))
```

Gradient Boosting MAE: 0.070872994978

Principal Component Analysis